


# Octree Scene Construction from a Depth Camera

Dave Kotfis and Patrick Cozzi

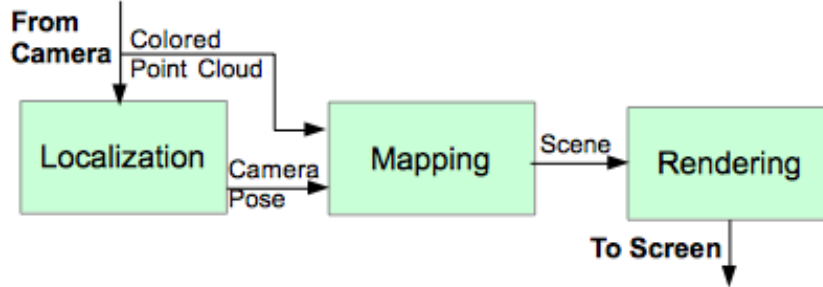
## 1.1 Overview

Data parallel GPU computing allows for real-time 3D reconstruction of scenes from depth cameras such as the Kinect sensor. Noise in the cameras depth measurements can be filtered over multiple image frames by representing the scene as a voxel-based map rather than as a collection of raw point clouds. However, a dense voxel grid representation is not suitable for large scenes or live rendering.

In this chapter, we present our methods that use CUDA to reconstruct



**Figure 1.1.** Augmented Reality: rendering a textured Stanford Bunny sitting on an office chair in our reconstructed 3D scene.



**Figure 1.2.** Top level system view of scene construction and rendering.

3D scenes from depth cameras at near real-time speeds. A scene is represented by a **sparse voxel octree (SVO)** structure that scales to large volumes. We render these scenes with CUDA and OpenGL using methods that eliminate the slow process of generating meshes from point clouds or voxel grids. We will describe an SVO representation of a scene and data parallel methods to update and expand from incrementally received colored point clouds.

### 1.1.1 Augmented Reality

In recent years, low cost depth cameras using structured light or time of flight methods have become common-place. These **RGB-D** (color + depth) cameras directly measure additional 3D information that previously could only be generated through sophisticated computer vision algorithms in software. These cameras are useful for creating models for 3D printing, computer vision for robotics, and creating immersive and interactive video game experiences.

**Augmented reality (AR)** is a field that lives on the boundary between computer graphics and vision to create experiences that blend artificial graphics with the real world. AR systems today typically render virtual objects in front of a raw depth camera frame. Future AR systems will seamlessly render virtual graphics blended with a live camera scene. Real scenes could be viewed with artificial lighting conditions, and with virtual objects that cast shadows. For the artificial objects to be rendered with consistent shading from arbitrary perspectives, a 3D scene needs to be constructed from the camera frames.

### 1.1.2 Localization

To reconstruct a scene, the movement of the camera between each frame must be determined so the points in each frame can be spatially correlated. GPU computing enables dense camera pose tracking techniques that match every pixel in 640x480 frames at 30 frames per second to track the motion of the camera without the need for a motion capture system. Previously, sparse techniques required detection of a smaller set of invariant features to track, which are not always available [Dryanovski et al. 13].

RGB-D cameras provide enough information to provide 3D positions and surface normals. The **iterative closest point (ICP)** algorithm attempts to align one frame to the previous by iteratively reducing the error between the points of each frame and the surfaces of the scene. Visual odometry with depth is a similar process that minimizes a photometric (color) error term rather than a geometric one [Steinbrucker et al. 11]. In different scenes, either geometric or photometric detail may be more prominent, so recent approaches use a combined error function that mixes the two [Whelan et al. 12].

The hard part is computing the error gradient fast enough to keep up with the 20+ fps frame rate while executing enough iterations (often 20+) for the solution to converge. If that rate cannot be maintained and frames are skipped, the space of possible transformations that must be searched to align the frames grows. This increases the computational burden, slowing the computation down even further and creating a vicious cycle that makes the process fail. GPU computing that exploits the parallelism of the computation is critical to achieve the speeds required to avoid this downward spiral.

### 1.1.3 Mapping

Reconstructing a scene requires a map representation to incrementally update and store data from each camera frame. There are many possible representations to do this, the simplest of which would be to concatenate each new point cloud by transforming all points according to the pose of the camera, assuming it is known. However, this size of this map would grow linearly with time, even when observing the same part of the scene, so it is not a suitable candidate for concurrent rendering. A standard RGB-D camera can generate several GB of raw data within only a minute.

A solution to the memory explosion is to filter the points in the map to truncate points to fall into 3D bins and remove duplicate points in space. This ensures that repeated observation of the same part of the scene will converge to a finite amount of memory over time. However, this will result in loss of detail, and the map will contain any noise produced by the camera

data. While the binning of the points is completely data parallel, the removal of point duplicates requires parallel sorting and reduction.

## 1.2 Previous Work and Limitations

### 1.2.1 KinectFusion

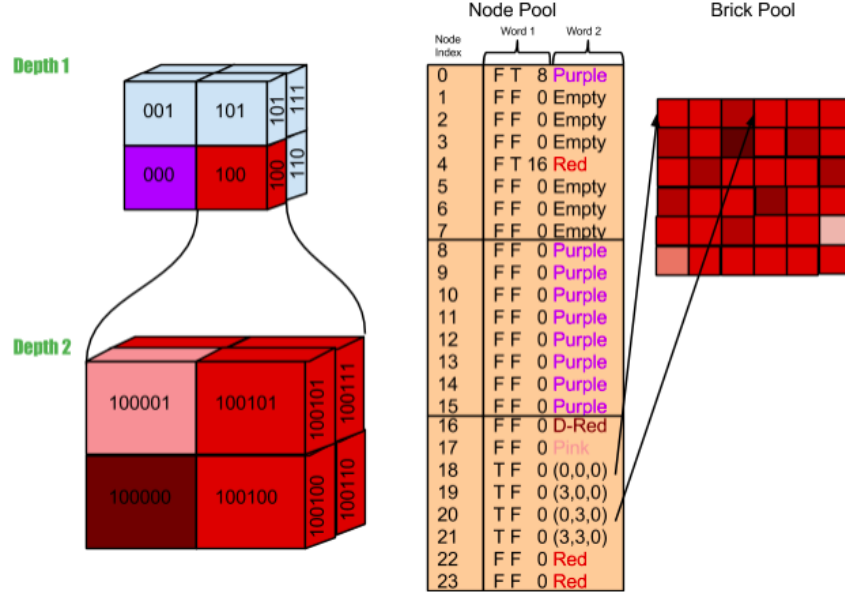
KinectFusion is a 3D reconstruction technique that attempts to filter the noise of incoming depth data by representing the map as a 3D voxel grid with a truncated signed distance function (TSDF) storing the distance from a surface [Newcombe et al. 11]. The values are truncated to avoid unnecessary computations in free space as well as reduce the amount of data required for surface representation. Building this grid is far more maintainable than storing a raw point cloud for each frame, as the redundancy both enables the sensor noise to be smoothed. It also avoids storing significant amounts of duplicate data, and is highly data parallel for GPU acceleration.

However, the memory footprint of a voxel grid approach scales poorly to large volumes. The dense representation requires voxel cells allocated in memory for the large amount of free space that will almost always be prominent in scenes. Also, while the voxel grid and TSDF are an appropriate representation for the surface function, it is inefficient for any color data. The rendering process either requires ray marching to directly render the grid, or a slow surface extraction and remeshing process, neither suitable for concurrent real-time rendering.

### 1.2.2 OctoMap

OctoMap is a probabilistic framework where the log-odds of occupancy are stored in an octree data structure [Hornung et al. 13]. The sparse octree structure overcomes the scalability limitations of a dense voxel grid by leaving free space unallocated in memory. OctoMap also filters sensor noise by assigning a probability of hit and miss that represent the noise of the sensor. Nodes in the tree are updated by logging each point from a point cloud as a hit. All points along the ray from the camera position to the end point are logged as a miss. This process takes place serially on a CPU, looping over each point in each frame.

The OctoMap is rendered by iterating through the leaves of the tree, and extracting cells that have a probability greater than 0.5 of being occupied. These voxels are rendered as cubes with edge length determined by the depth of the corresponding node in the octree. This framework is most commonly used with LIDAR sensors, which have only a few points per



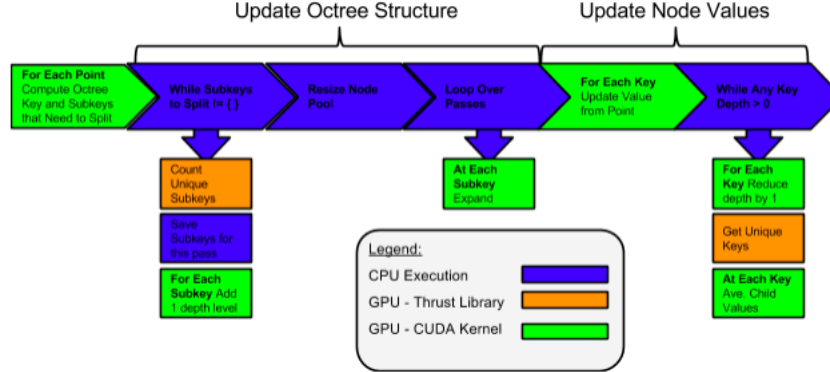
**Figure 1.3.** Sparse Voxel Octree data structure in linear GPU memory. Utilizes keys based on Morton Codes to uniquely index nodes. Compact structure uses 64 bits per node. For hardware interpolation of values within the tree, node values can be backed by a brick in texture memory.

scan which has little benefit from parallelization. An RGB-D sensor would provide millions of points per frame which could be parallelized. However, the pointer-based octree structure that it uses is less suitable for GPU parallelization than a stackless linear octree.

## 1.3 Octree Scene Representation

### 1.3.1 Data Format

We have developed a sparse octree representation of a scene on a GPU, along with methods to efficiently update and expand it from incrementally received colored point clouds. The GPU data structure is based on the work of GigaVoxels [Cra 09] that uses a Node Pool in linear memory and a Brick Pool in texture memory. The nodes are composed of two 32-bit words. The first word has two single bit flags, and 30 bits for the index of the first child node. The second word either holds an RGBA value, or the location in the brick pool to be used when interpolating values within the



**Figure 1.4.** Program flow for updating the sparse octree from a point cloud frame. The process of updating the octree map from each incoming point cloud starts by counting how many new octree nodes must be created and resizing the node pool. Then we can filter the updated color values through the tree.

node.

Although the sparse octree does not allocate every node of the tree in memory, Morton Codes are used for unique identification of voxels. Here is an example code: 1 001 010 111. The code starts with a leading 1 to identify the length of the code, and thus the depth in the tree. After that, the code is made up of a series of 3-bit tuples that indicate a high or low value on the binary split of the x, y, and z dimensions respectively.

Using a 32-bit integer, this can represent 10 levels of depth in the tree. However, this is insufficient for mapping with a Kinect camera. The Kinect has a range of 1-10 meters in depth resolution, and doing a back of the envelope calculation for the horizontal resolution (480 pixels, 5m range, 43 degree field of view) shows that the camera will typically provide sub-centimeter resolution. A 10 meter edge volume can only achieve 1 cm resolution using 10 levels of depth. Therefore, we have transitioned to representing these keys with long integers (64-bit), which could represent more than kilometers of volume at millimeter precision, if needed. Figure 1.3 and Listing 1.1 provide descriptions of our data format.

```

struct char4 {
    char x, y, z, w;
};

struct Octree {
    //The node data in GPU memory.
    //Each node is 2 unsigned int s long.

```

```

unsigned int* node_pool;
//The number of nodes allocated in the node pool
int size;
//The half length of each edge of the root node of the octree
float edge_length;
//The 3D position of the center of the root node of the octree
glm::vec3 center;
//The brick pool data in CUDA texture memory.
//Note:Our examples are limited to use of the node pool only.
cudaArray* brick_pool;
};

struct PointCloud {
//The 3D position of each point in the point cloud
glm::vec3* positions;
//The corresponding RGBA color of each corresponding point
char4* colors;
//The number of points in the cloud
int size;
};

```

**Listing 1.1.** Data structures representing a sparse linear octree and colored point cloud data in GPU memory.

### 1.3.2 Updating the Map

Because our data structure is sparse, each new point cloud frame may contain points in parts of space that were previously unallocated in memory. For this reason, updating the map requires two steps: resizing the octree structure into newly observed space, and updating the color values within the tree with the new observations. Figure 1.4 shows the program flow for updating the map in more detail.

**Update Octree Structure** To expand our scene into unallocated space, we first must determine which new points correspond to unallocated nodes. We do this by computing the Morton Code for each point to determine its location in the octree. Fortunately, we can do this with only the constant octree parameters, its size and center location, without the need for any data within the octree. This makes the calculation completely data parallel over the incoming point cloud positions. The process of computing Morton Codes is in Listing 1.2

```

typedef long long int octkey;

__device__ octkey computeKey(const glm::vec3& point,
    glm::vec3 center, const int tree_depth,
    float edge_length) {
    //Initialize the output value with a leading 1

```

```

//to specify the depth
octkey morton = 1;

for (int i = 0; i < tree_depth; i++) {
    morton = morton << 3;

    //Determine which octant the point lies in
    uint8_t x = point.x > center.x;
    uint8_t y = point.y > center.y;
    uint8_t z = point.z > center.z;

    //Update the code
    morton += (x + 2*y + 4*z);

    //Update the edge length
    edge_length /= 2.0f;

    //Update the center
    center.x += edge_length * (x ? 1 : -1);
    center.y += edge_length * (y ? 1 : -1);
    center.z += edge_length * (z ? 1 : -1);
}
return morton;
}

```

**Listing 1.2.** CUDA device function to compute a Morton Code for a point. A kernel that parallelizes over points should call this.

TODO: Describe the rest of the update svo structure process here.

**Update Node Values** To update the node values, we use the alpha channel of RGBA to encode a pseudo log probability of hit for each cell. When allocated, we initialize our cells to  $\alpha = 127$ . For a Kinect sensor, we use a probability of hit such that each observation adds 2 to the alpha value. This means that the more often we observe a point within a portion of space, the more confident that we are that the node is occupied. This helps to filter sensor noise in depth measurements by ensuring that we consistently receive point returns from a location before considering it to be filled in.

We also filter the color values received by the camera by using a running average, using the alpha channel as a weight function. Listing 1.3 shows the update and filtering process for each node. After the values are updated in the leaves of the octree, we can trickle them into the inner limbs of the tree by having each parent assume a value that averages their children.

```

--device-- int getFirstValueAndShiftDown(octkey& key) {
    int depth = depthFromKey(key);
    int value = getValueFromKey(key, depth-1);
    key -= ((8 + value) << 3 * (depth - 1));
    key += (1 << 3 * (depth - 1));
}

```



```

    return value;
}

--global-- void fillNodes(const octkey* keys, int numKeys,
                        const char4* values, unsigned int* octree_data) {

    int index = blockIdx.x * blockDim.x + threadIdx.x;

    //Don't do anything if out of bounds
    if (index >= numKeys) {
        return;
    }

    //Get the key for this thread
    octkey key = keys[index];

    //Check for invalid key
    if (key == 1) {
        return;
    }

    int node_idx = 0;
    int child_idx = 0;
    while (key != 1) {
        //Get the child number from the first three bits of the
        //morton code
        node_idx = child_idx + getFirstValueAndShiftDown(key);

        if (!octree_data[2 * node_idx] & 0x40000000) {
            return;
        }

        //The lowest 30 bits are the address of the child nodes
        child_idx = octree_data[2 * node_idx] & 0x3FFFFFFF;
    }

    char4 new_value = values[index];
    unsigned int current_value = octree_data[2 * node_idx + 1];

    char4 current;
    short current_alpha = current_value >> 24;
    current.r = current_value & 0xFF;
    current.g = (current_value >> 8) & 0xFF;
    current.b = (current_value >> 16) & 0xFF;

    //Implement a pseudo low-pass filter with laplace smoothing
    float f1 = (1 - ((float)current_alpha/256.0f));
    float f2 = (float)current_alpha / 256.0f;
    new_value.r = new_value.r * f1 + current.r * f2;
    new_value.g = new_value.g * f1 + current.g * f2;
    new_value.b = new_value.b * f1 + current.b * f2;
    octree_data[2 * node_idx + 1] = ((int)new_value.r) +
        ((int)new_value.g << 8) + ((int)new_value.b << 16) +
        (min(255, current_alpha + 2) << 24);
}

```

```
}

```

**Listing 1.3.** CUDA kernel for updating values stored in octree nodes based on newly observed colors.

### 1.3.3 Dynamic Scenes

When building a scene where all objects are static, it would be sufficient to update the map in only an additive fashion as discussed earlier. However, when objects are moving it becomes necessary to have an update process that can remove parts of the map when they are observed to be unoccupied. Similar to OctoMap, we do this by processing the free space between the camera origin and each point in our point cloud. In each update, these nodes are observed to be free. Rather than adding an additional registered hit to these nodes, we register them as misses. With enough misses, these nodes will eventually return to being unoccupied.

Once these nodes are completely unoccupied ( $\alpha = 0$ ), the memory for them is released. Rather than the expensive process of shifting all of the data in memory to fill in these holes, maintaining a list of free memory slots allows future tree expansions to fill data into them first.

### 1.3.4 Managing Memory

The sparse octree used to represent a reconstructed 3D map will quickly grow too large to fit entirely in GPU memory. Reconstructing a typical office room at 1 cm resolution will often take as much as 6-8 GB.

To handle this, we developed an out-of-core memory management framework for the octree. At first glance, this framework is a standard stack-based octree on the CPU. However, each node in the tree has an additional boolean flag indicating whether the node is a subtree that is located in linear GPU memory. It also holds a pointer to its location on the GPU as well as its size.

Next, these nodes can push/pull the data to and from the GPU. The push method uses recursion to convert the stack-based data into a linear array in CPU memory, then copies the memory to the GPU. It avoids the need to over allocate or reallocate the size of the linear memory by first recursing through the node's children to determine the size of the subtree. The pull method copies the linear memory back to the CPU, then uses it to recursively generate it as a stack-based structure.

We use a Least Recently Used (LRU) approach where all methods operating on the tree must provide an associated bounding box of the area that they will affect. First, this allows us to make sure that the entire affected volume is currently on the GPU before attempting to perform the operation. The octree will also keep a history of the N most recently used



**Figure 1.5.** Octree scene constructed from a live Kinect camera stream using CUDA. Rendering is performed with OpenGL instancing of a cube with a texture buffer object containing voxel centers, colors, and scales.

bounding boxes. When space needs to be freed, it will take the union of these stored bounding boxes and pull data that lies outside of this region back to the CPU.

## 1.4 Rendering Techniques

### 1.4.1 Extracting and Instancing Voxel Cubes

The brute force method for rendering the SVO map is to extract the color values and 3D positions of each occupied leaf node. With these values, we can render a cube at each center position with a scale based on the depth in the SVO.

Extracting the voxels requires two steps. First, in a prepass where each CUDA thread is assigned a Morton Code each traverses into the SVO to determine whether the node with the corresponding code is occupied. We start with a set of keys at the minimum depth, and iteratively create the 8 child keys for the occupied nodes, and remove the unoccupied node keys.



**Figure 1.6.** Multiple renders of the same scene, both with voxel cone tracing. On the left, the max resolution is 1 cm, while on the right it is capped at 16 cm.

Once we have determined the valid keys, we allocate space for our resulting data, and extract it from the SVO into the buffer. We decode the Morton Codes back into 3D positions for each voxel.

Once we have the position and color for each occupied voxel, we map it to an OpenGL **Texture Buffer Object (TBO)** which is used by our vertex shader that instances a colored cube to represent them (Listing 1.4).

### 1.4.2 Voxel Cone Tracing

Voxel Cone Tracing is a physically based rendering technique similar to ray tracing [Crassin et al. 11]. It exploits the SVO data structure to avoid Monte Carlo integration of multiple rays to approximate the integral of the rendering equation. Instead, it approximates a cone by sampling values at higher levels of the SVO as the cone becomes wider. If all of the needed lighting information is incorporated into the octree, mip-mapping the values into the inner tree branches and texture interpolation performs the integration step inherently.

We used Voxel Cone Tracing to render our scene with CUDA. For each pixel, a CUDA thread traverses along a ray and samples a value from the SVO. It continues to integrate its total color value using the alpha channel until it reaches its max of 255, or until the ray reaches a max length (usually 10 m).

```
#version 420

uniform mat4 u_mvMatrix;
uniform mat3 u_normMatrix;
uniform float u_scale;

out vec3 fs_position;
```

```

out vec3 fs_normal;
out vec3 fs_color;

layout (location = 0) in vec4 vox_cent;
layout (location = 1) in vec4 vox_color;

layout (binding = 0) uniform samplerBuffer voxel_centers;
layout (binding = 1) uniform samplerBuffer voxel_colors;

const vec3 cube_vert[8] = vec3[8](
    vec3(-1.0, -1.0, 1.0),
    vec3(1.0, -1.0, 1.0),
    vec3(1.0, 1.0, 1.0),
    vec3(-1.0, 1.0, 1.0),
    vec3(-1.0, -1.0, -1.0),
    vec3(1.0, -1.0, -1.0),
    vec3(1.0, 1.0, -1.0),
    vec3(-1.0, 1.0, -1.0)
);

const int cube_ind[36] = int[36] (
    0, 1, 2, 2, 3, 0,
    3, 2, 6, 6, 7, 3,
    7, 6, 5, 5, 4, 7,
    4, 0, 3, 3, 7, 4,
    0, 1, 5, 5, 4, 0,
    1, 5, 6, 6, 2, 1
);

void main (void){
    gl_Position = u_mvMatrix *
        vec4(cube_vert[cube_ind[gl_VertexID]]*u_scale +
            vec3(texelFetch(voxel_centers, gl_InstanceID)), 1.0);
    fs_position = gl_Position.xyz;
    fs_normal = u_normMatrix *
        normalize(cube_vert[cube_ind[gl_VertexID]]);
    fs_color = vec3(texelFetch(voxel_colors, gl_InstanceID));
}


```

**Listing 1.4.** GLSL vertex shader for instancing of colored voxel cubes using a TBO bound from CUDA.

## 1.5 Results

We test the time required to expand, update, and filter an SVO scene with an updated point cloud frame from a Kinect sensor. We found the time to increase logarithmically with the number of allocated nodes in the SVO (Figure 1.7). The kernels that update the SVO execute serially in tree depth, but parallel over the nodes in each depth. The octree structure divides the nodes so that we can expect the depth to increase logarithmically

with the number of nodes.



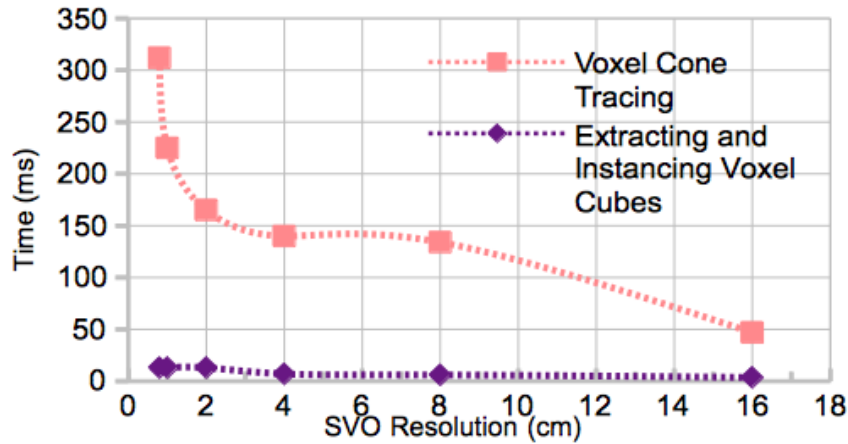
**Figure 1.7.** Evaluation of updating the SVO scene from a Kinect camera using an NVidia GTX 770 with 2 GB memory. The same scene is updated with multiple maximum depths. The edge length of the full SVO is 1.96 meters. We evaluate the update time, and compare it with the change in the number of allocated nodes in the octree.

We compare the rendering time between both the voxel instancing and voxel cone tracing approaches with an identical scene at multiple levels of resolution. We found that the voxel instancing approach has steady real-time performance at all resolutions tested. Even at the lowest resolution, the voxel cone tracing technique was not real-time. The run time for VCT grows exponentially as the resolution increases (Figure 1.8).

## 1.6 Conclusion and Future Work

We have found that use of an SVO map allows for memory efficient scene construction. Camera noise is quickly filtered out within a few frames to create stable scenes. For debug views, voxel extraction and instanced rendering is useful for rendering values of the map at different levels of resolution. However, voxel cone tracing requires minimal additional computational cost and can render the scene at different views with similar quality to that of the original.

We would like to explore use of intrinsic images work in preprocessing the color values before adding them to the map. This would allow us to re-cast an artificial light into the scene without the rendering artifacts that



**Figure 1.8.** Here we render the SVO scene with both voxel extraction and instancing and cone tracing with the same scene as the previous figure. Voxel extraction and instancing achieves real-time performance at every resolution tested, but cone tracing slows down below real-time higher resolutions than 16 cm.

we expect from improper shading. Rendering with a virtual light source would also blend virtual objects into the scene by casting shadows.

Also, today we are only able to add static virtual objects to our constructed scenes. It would be useful for dynamic virtual objects to move efficiently within the SVO.

## Bibliography

- [Cra 09] *GigaVoxels: ray-guided streaming for efficient and detailed voxel rendering*, 2009.
- [Crassin et al. 11] Cyril Crassin, Fabrice Neyret, Miguel Sainz, Simon Green, and Elmar Eisemann. “Interactive Indirect Illumination Using Voxel Cone Tracing.” *Computer Graphics Forum*.
- [Dryanovski et al. 13] Ivan Dryanovski, Roberto G. Valenti, and Jizhong Xiao. “Fast Visual Odometry and Mapping from RGB-D Data.” In *IEEE International Conference on Robotics and Automation (ICRA)*, 2013.
- [Hornung et al. 13] Armin Hornung, Kai M. Wurm, Maren Bennewitz, Cyril Stachniss, and Wolfram Burgard. “OctoMap: An Efficient Prob-

abilistic 3D Mapping Framework Based on Octrees.” *Autonomous Robots*.

- [Newcombe et al. 11] R.A. Newcombe, S. Izadi, O. Hilliges, D. Molyneaux, D. Kim, A. J. Davison, P. Kohli, J. Shotton, S. Hodges, and A. Fitzgibbon. “KinectFusion: Real-time Dense Surface Mapping and Tracking.” In *IEEE Int. Symposium on Mixed and Augmented Reality (ISMAR)*, 2011.
- [Steinbrucker et al. 11] F. Steinbrucker, J. Sturm, and D. Cremers. “Real-time visual odometry from dense RGB-D images.” In *ICCV Workshop on Live Dense Reconstruction with Moving Cameras*, 2011.
- [Whelan et al. 12] T. Whelan, J. McDonald, M. Fallon M. Kaess, H. Johannsson, and J. Leonard. “Kintinuous: Spatially Extended KinectFusion.” In *RSS Workshop on RGB-D: Advanced Reasoning with Depth Cameras*, 2012.